

An introduction to Bayesian modelling with JAGS and R

Francisco Rodriguez-Sanchez (@frod_san)

April 2015

This is a practical workshop

BUT do read the literature e.g.

- [Data analysis using regression and multilevel/hierarchical models](#)
- [Bayesian data analysis](#)
- [Bayesian methods for ecology](#)
- [The BUGS book](#)
- [Introduction to WinBUGS for ecologists](#)
- [Models for ecological data](#)
- and many more

Bayesian modelling software

- WinBUGS/OpenBUGS
- JAGS
- STAN
- Filzbach
- Nimble
- Many R packages: MCMCpack, MCMCglmm, LaplacesDemon, r-inla, etc (see [Bayesian task view](#))

Why JAGS?



- Very similar to BUGS, both very popular
- Gate to other software e.g. STAN, Filzbach, etc
- Easy to start, can deal with complex models too (open-ended modelling)
- But look for specific implementations of your analysis (e.g. hSDM)

Why R?

Just kidding

:)



The very basics: linear regr



$$y_i = a + bx_i + \epsilon_i$$

How many parameters?

The very basics: linear regr



$$y_i = a + bx_i + \epsilon_i$$

$$\epsilon \sim N(0, \sigma^2)$$

Or also

$$y_i \sim N(\mu_i, \sigma^2)$$

$$\mu_i = a + bx_i$$


Our dataset: tree heights and DBH

- One species
- 10 plots
- 1000 trees
- Number of trees per plot ranging from 4 to 392

```
trees <- read.csv("trees.csv")  
summary(trees[,1:3])
```

plot	dbh	height
Min. : 1.0	Min. : 5.06	Min. :13.40
1st Qu.: 1.0	1st Qu.:17.69	1st Qu.:29.68
Median : 2.0	Median :28.62	Median :36.55
Mean : 2.7	Mean :27.88	Mean :36.51
3rd Qu.: 4.0	3rd Qu.:38.97	3rd Qu.:43.33
Max. :10.0	Max. :49.92	Max. :59.30

What's the relationship between DBH and height?



First step: linear regression (lm)

```
simple.lm <- lm(height ~ dbh, data=trees)
arm::display(simple.lm)
```

```
lm(formula = height ~ dbh, data = trees)
      coef.est coef.se
(Intercept) 19.34    0.31
dbh          0.62    0.01
---
n = 1000, k = 2
residual sd = 4.09, R-Squared = 0.79
```

Interpretation?

Always centre continuous variables



```
summary(trees$dbh)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
5.06	17.69	28.62	27.88	38.96	49.92

```
trees$dbh.c <- trees$dbh - 25
```

So, all parameters will be referred to a 25 cm DBH tree.

Linear regression with cen



```
lm(formula = height ~ dbh.c, data = trees)
```

	coef.est	coef.se
(Intercept)	34.73	0.13
dbh.c	0.62	0.01

n = 1000, k = 2

residual sd = 4.09, R-Squared = 0.79

Let's make it Bayesian

Things we'll need

- Data
- A function describing the model (including **priors**)
- Decide number of MCMC chains
- Define initial values
- Decide number of iterations (and burnin)
- Choose parameters to save

Specify the model as an R function

```
modell1 <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)    # tau = precision (inverse of variance)  
    mu[i] <- alfa + beta*dbh[i]      # centred diameter  
  }  
  
  # PRIORS (vague or weakly informative)  
  alfa ~ dunif(1, 100)              # prior for average height of a 25-cm-DBH tree  
  beta ~ dunif(0, 10)               # how much do we expect height to scale with DBH?  
  tau <- pow(sigma, -2)              # tau = 1/sigma^2  
  sigma ~ dunif(0, 50)              # residual standard deviation  
}
```

A note on priors

Avoid 'non-informative' priors (see [this](#) and [this](#))


Use *weakly informative* (e.g. bounded Uniform, Normal with reasonable parameters, Cauchy...)

or *strongly informative* priors based on previous knowledge and common sense.

Example: estimating people height across countries


Unreasonable prior

```
plot(density(rnorm(1000, 0, 1000)),  
     main="", xlab="Height (m)")
```



Reasonable prior

```
plot(density(rnorm(1000, 2, 0.5)),  
     main="", xlab="Height (m)")
```



(from STAN manual)

Next step: create list with data

```
data <- list(height = trees$height,  
             dbhc = trees$dbh.c)
```

Now call JAGS to run the model

```
m1 <- jags(data,  
           model.file=model1,  
           parameters.to.save = c("alfa", "beta", "sigma"),  
           n.chains=3,  
           inits=NULL,  
           n.iter=10,  
           n.burnin=5)
```

Compiling model graph

Resolving undeclared variables

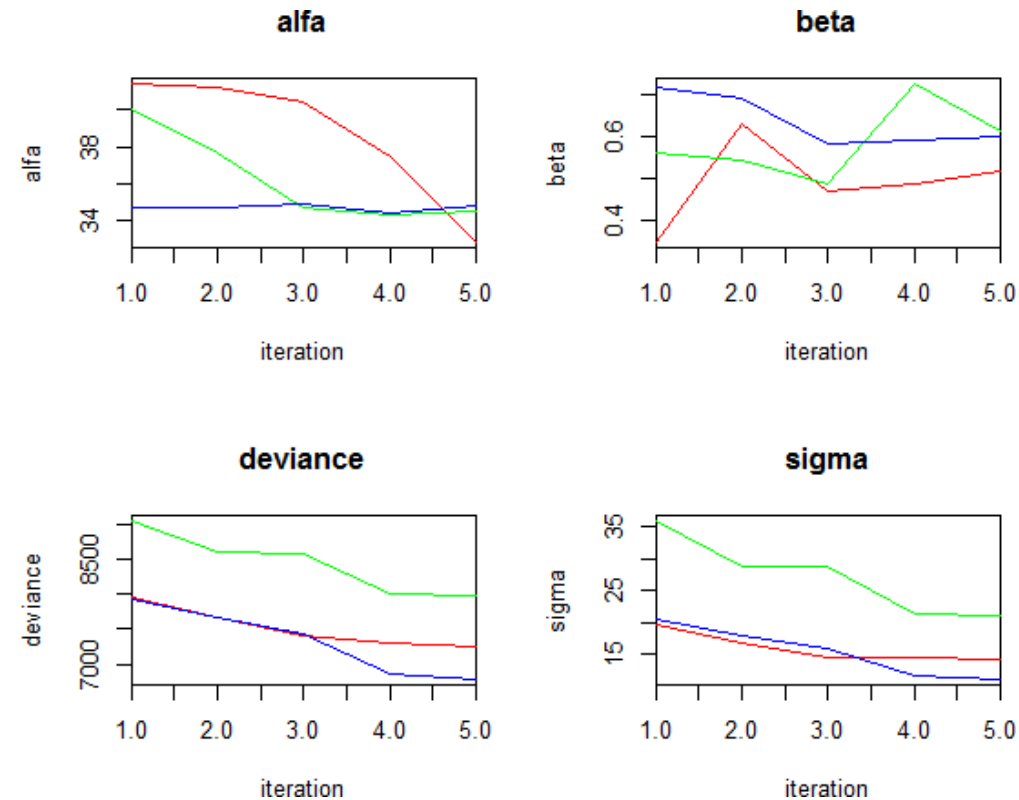
Allocating nodes

Graph Size: 3787

Initializing model

Viewing MCMC in action

```
traceplot(m1, ask=FALSE, mfrow=c(2,2))
```



Obviously we haven't achieved convergence yet...

Let's run JAGS for longer

```
m1 <- jags(data,  
           model.file=model1,  
           parameters.to.save = c("alfa", "beta", "sigma"),  
           n.chains=3,  
           inits=NULL,  
           n.iter=10000,  
           n.burnin=5000)
```

Compiling model graph

Resolving undeclared variables

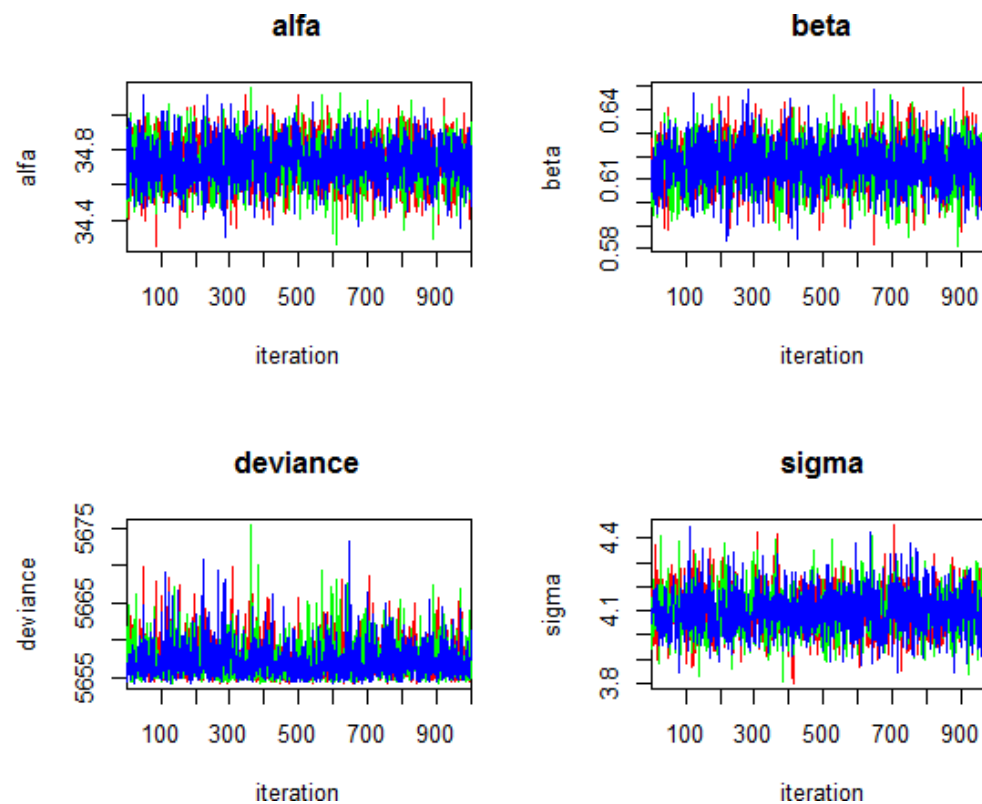
Allocating nodes

Graph Size: 3787

Initializing model

Traceplots

```
traceplot(m1, ask=FALSE, mfrow=c(2,2))
```



Results

```
Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpwpZdlg/model18fc32fa311f.txt", f
  3 chains, each with 10000 iterations (first 5000 discarded), n.thin = 5
n.sims = 3000 iterations saved
```

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alfa	34.729	0.134	34.464	34.980	1.001	3000
beta	0.616	0.010	0.596	0.636	1.002	1700
sigma	4.099	0.092	3.926	4.280	1.001	3000
deviance	5657.303	2.529	5654.453	5663.886	1.001	2300

For each parameter, n.eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor (at convergence, Rhat=1).

DIC info (using the rule, $pD = \text{var}(\text{deviance})/2$)

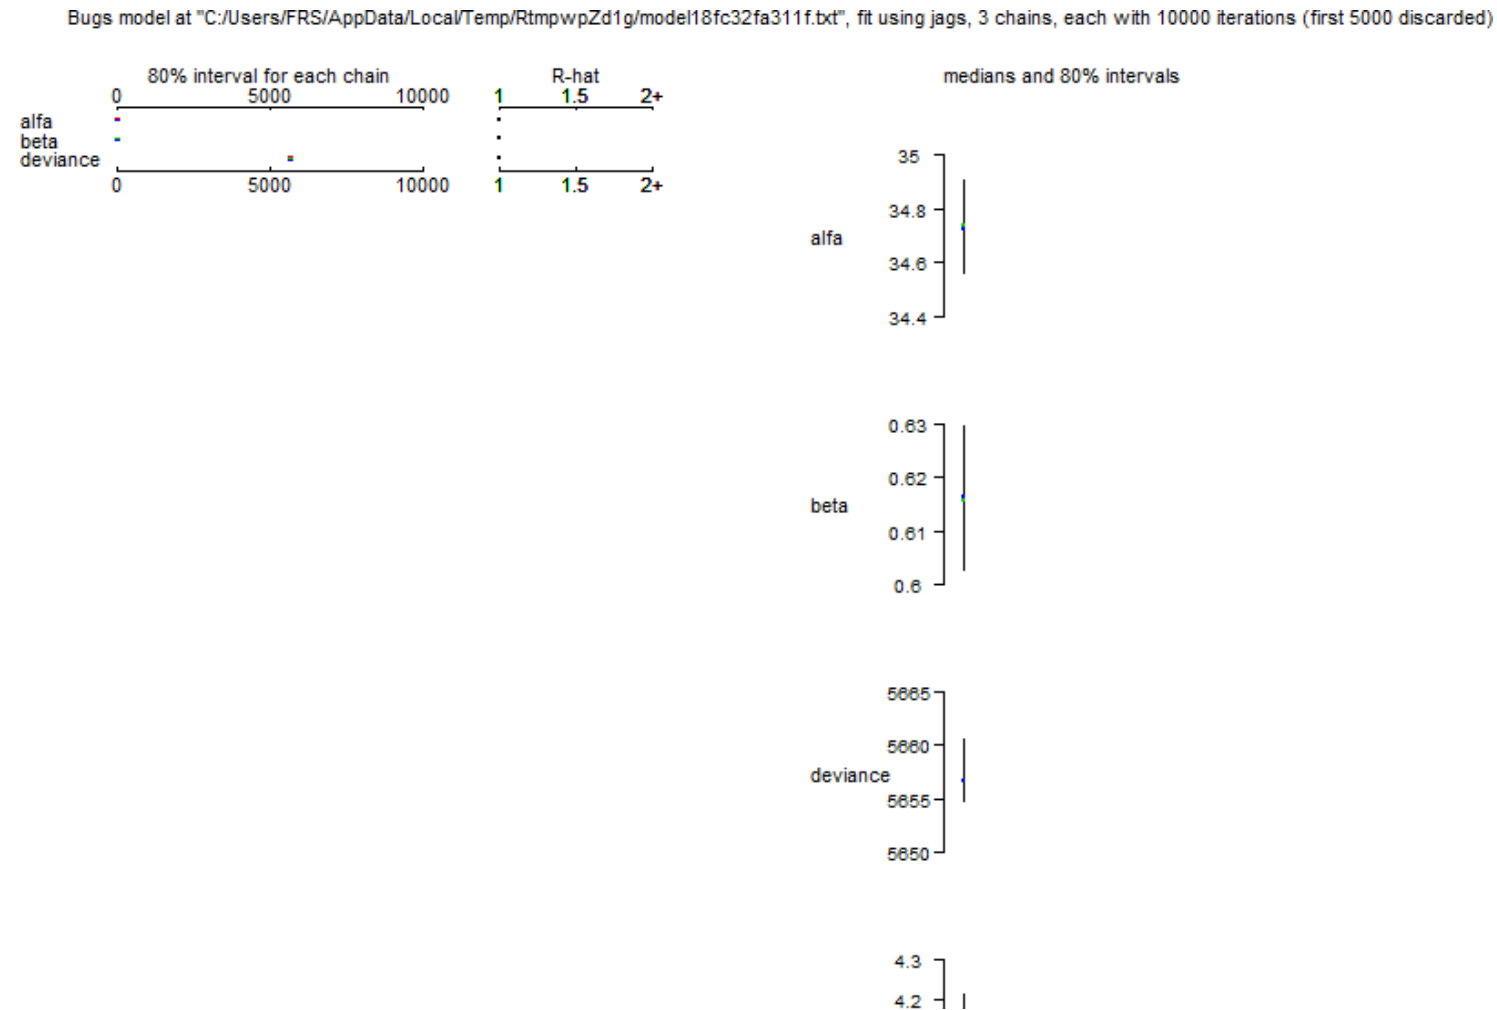
$pD = 3.2$ and $DIC = 5660.5$

DIC is an estimate of expected predictive error (lower deviance is better).

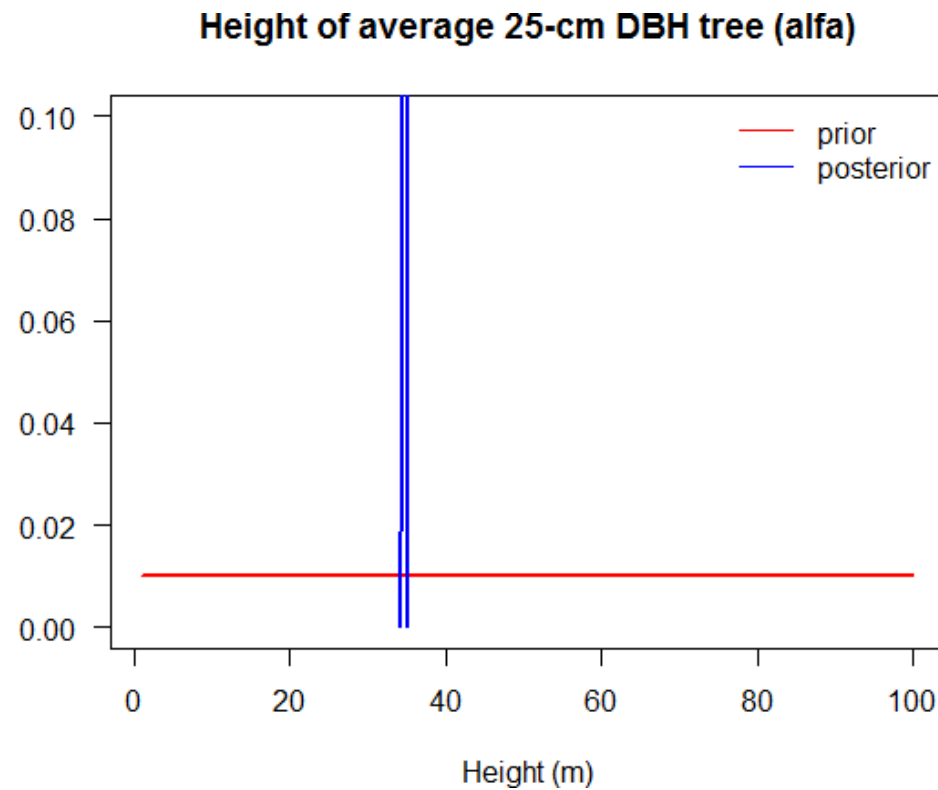
Results pretty similar to simple.lm (because of vague priors)

A plot of the whole model

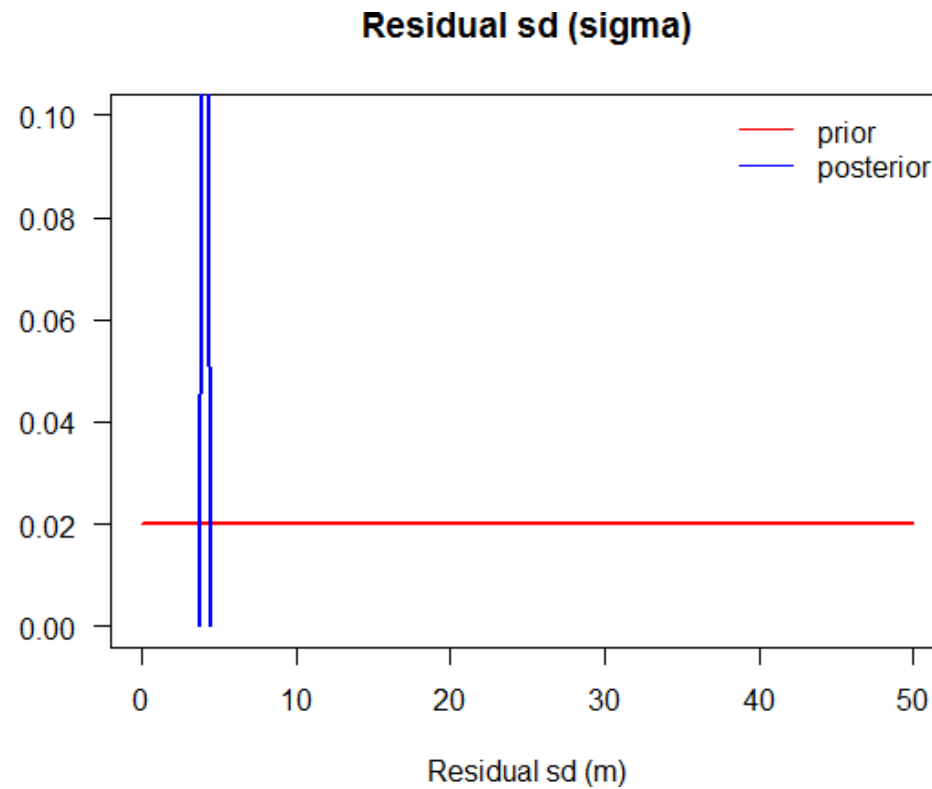
```
plot(m1)
```



Comparing prior and posterior densities



Comparing prior and posterior densities



Now using Normal vague priors

Model with Normal priors

```
modellb <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)    # tau = precision (inverse of variance)  
    mu[i] <- alfa + beta*dbh[i]      # centred diameter  
  }  
  
  # PRIORS  
  alfa ~ dnorm(0, 0.001)             # prior for intercept  
  beta ~ dnorm(0, 0.001)             # prior for beta (slope)  
  tau <- pow(sigma, -2)              # tau = 1/sigma^2  
  sigma ~ dunif(0, 50)               # residual standard deviation  
}
```

Calling JAGS

```
m1 <- jags(data,  
            model.file=model1b,  
            parameters.to.save = c("alfa", "beta", "sigma"),  
            n.chains=3,  
            inits=NULL,  
            n.iter=10000,  
            n.burnin=5000)
```

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 3785

Initializing model

Results

```
Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpwpZdlg/model18fcecc61de.txt", fi
  3 chains, each with 10000 iterations (first 5000 discarded), n.thin = 5
n.sims = 3000 iterations saved
```

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alfa	34.732	0.131	34.470	34.980	1.001	3000
beta	0.615	0.010	0.595	0.636	1.003	950
sigma	4.098	0.095	3.918	4.286	1.001	3000
deviance	5657.292	2.526	5654.481	5664.219	1.004	600

For each parameter, n.eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor (at convergence, Rhat=1).

DIC info (using the rule, $pD = \text{var}(\text{deviance})/2$)

$pD = 3.2$ and $DIC = 5660.5$

DIC is an estimate of expected predictive error (lower deviance is better).

Very similar


Bayesian inference

$$p(\theta | x) = \frac{p(x | \theta)p(\theta)}{p(x)}$$

Posterior distribution \propto Likelihood \times Prior distribution

Varying-intercept models

Accounting for plot effects



Do it yourself using `lm`

lm results

```
lm.plot <- lm(height ~ factor(plot) + dbh.c, data=trees)
```

```
lm(formula = height ~ factor(plot) + dbh.c, data = trees)
```

	coef.est	coef.se
(Intercept)	32.13	0.16
factor(plot)2	6.50	0.26
factor(plot)3	4.36	0.35
factor(plot)4	1.93	0.36
factor(plot)5	3.64	0.34
factor(plot)6	4.20	0.42
factor(plot)7	-0.18	0.67
factor(plot)8	-5.31	0.89
factor(plot)9	5.44	1.09
factor(plot)10	2.26	1.37
dbh.c	0.62	0.01


```
n = 1000, k = 11
```

```
residual sd = 3.04, R-Squared = 0.88
```


Interpretation?

Single vs varying intercept

Pooling all plots



Different intercept for each plot



Let's make it Bayesian

Things we'll need

- Data
- A function describing the model (including **priors**)
- number of MCMC chains
- initial values
- number of iterations (and burnin)
- parameters to save

Bayesian varying-intercept model with no pooling

```
model2 <- function(){  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)    # tau = precision (inverse of variance)  
    mu[i] <- alfa[plot[i]] + beta*dbh[i]    # centred diameter  
  }  
  # PRIORS  
  #alfa ~ dnorm(0, .001)  
  for (j in 1:10){  
    alfa[j] ~ dnorm(0, .001) # Plot effects drawn from Normal distribution  
                           # with large **fixed** variance  
  }  
  beta ~ dnorm(0, .001)  
  tau <- pow(sigma, -2)      # tau = 1/sigma^2  
  sigma ~ dunif(0, 50)  
}
```

This fits same model as `lm.plot`

Call JAGS

```
data <- list(height=trees$height,  
             dbhc=trees$dbh.c,  
             plot=trees$plot)  
m2 <- jags(data,  
           model.file=model2,  
           parameters.to.save = c("alfa", "beta", "sigma"),  
           n.chains=3,  
           inits=NULL,  
           n.iter=10000,  
           n.burnin=5000)
```

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 4880

Initializing model

Results

Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpwpZdlg/model18fc4f7b76af.txt", f
 3 chains, each with 10000 iterations (first 5000 discarded), n.thin = 5
 n.sims = 3000 iterations saved

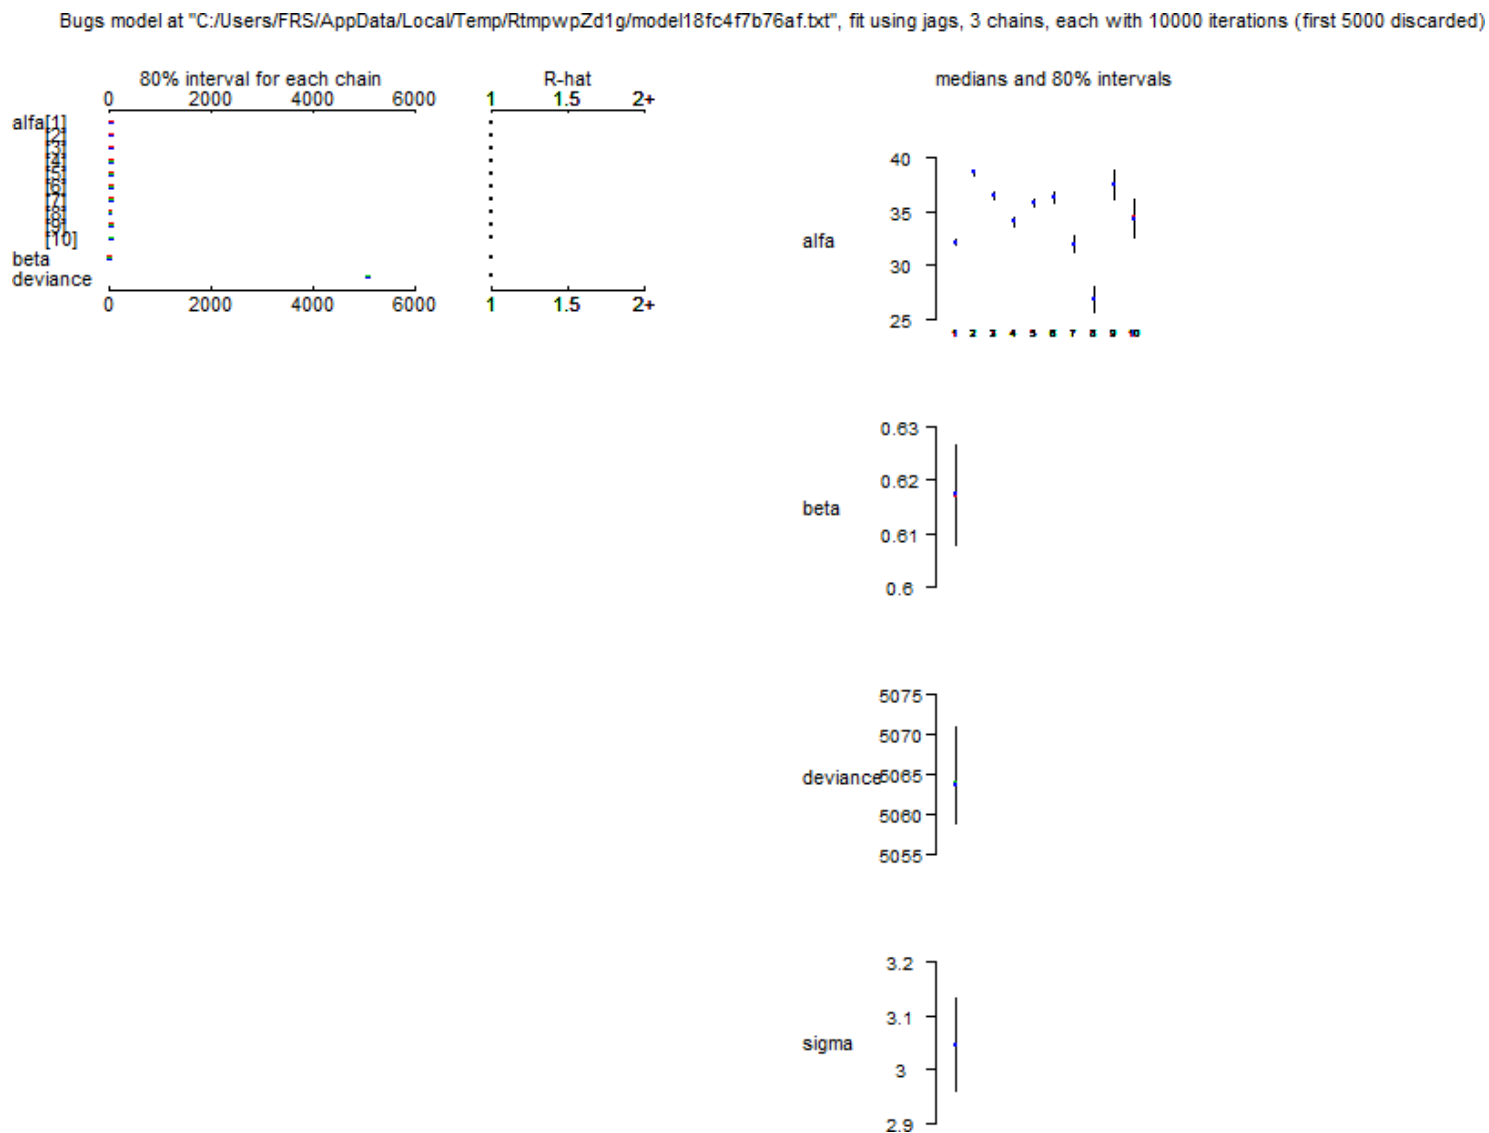
	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alfa[1]	32.126	0.152	31.828	32.420	1.001	3000
alfa[2]	38.628	0.208	38.219	39.037	1.003	750
alfa[3]	36.472	0.315	35.871	37.085	1.001	2700
alfa[4]	34.058	0.319	33.425	34.683	1.002	1000
alfa[5]	35.760	0.310	35.137	36.372	1.002	1500
alfa[6]	36.314	0.399	35.571	37.082	1.001	3000
alfa[7]	31.949	0.640	30.680	33.210	1.002	1800
alfa[8]	26.790	0.870	25.086	28.452	1.001	3000
alfa[9]	37.498	1.057	35.490	39.596	1.001	3000
alfa[10]	34.322	1.384	31.648	37.016	1.002	1900
beta	0.617	0.007	0.602	0.632	1.001	3000
sigma	3.047	0.068	2.915	3.188	1.001	3000
deviance	5064.492	4.914	5056.952	5075.810	1.001	3000

For each parameter, n.eff is a crude measure of effective sample size,
 and Rhat is the potential scale reduction factor (at convergence, Rhat=1).

38/73

DIC info (using the rule. $\text{pD} = \text{var}(\text{deviance})/2$)

Plot whole model

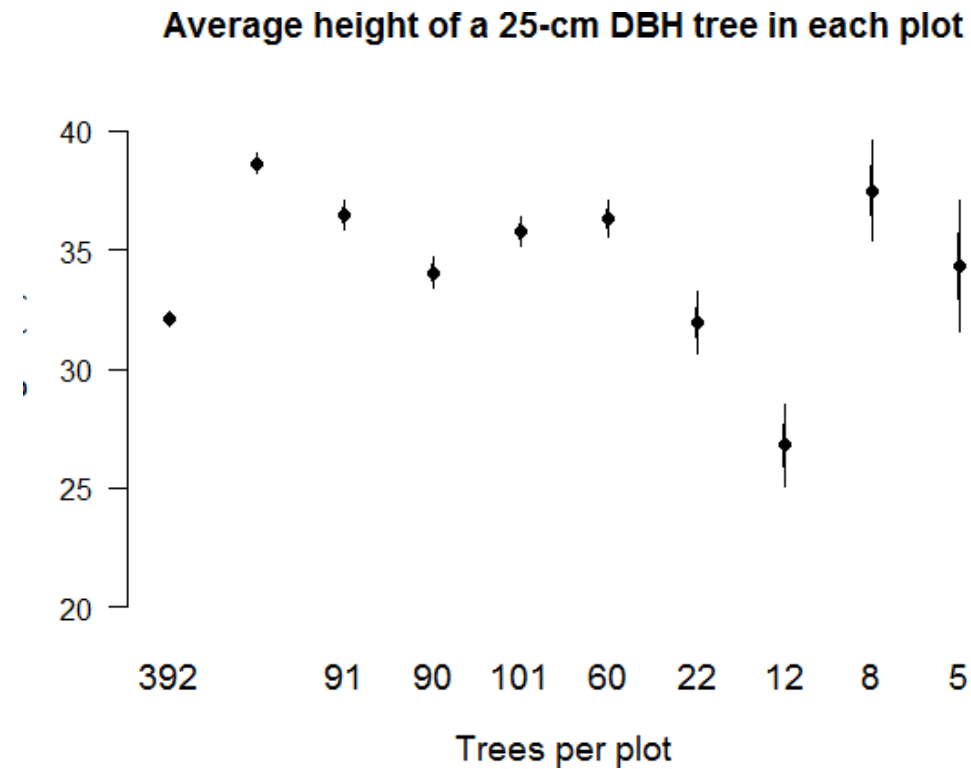


The varying-intercept model is much better

$DIC(m1) = 5660$

$DIC(m2) = 5077$

Estimation of plot effects improves with sample size



Varying-intercepts with pooling

(mixed/multilevel/hierarchical model)

Multilevel model with varying intercepts

$$y_i = a_j + bx_i + \varepsilon_i$$

$$a_j \sim N(0, \tau^2)$$

$$\varepsilon_i \sim N(0, \sigma^2)$$

In our example:

$$Height_i = plot_j + bDBH_i + \varepsilon_i$$

$$plot_j \sim N(0, \tau^2)$$

$$\varepsilon_i \sim N(0, \sigma^2)$$

Fitting mixed models with lmer

```
mixed <- lmer(height ~ dbh.c + (1|plot), data = trees)
```

```
lmer(formula = height ~ dbh.c + (1 | plot), data = trees)
```

	coef.est	coef.se
(Intercept)	34.43	1.08
dbh.c	0.62	0.01

Error terms:

Groups	Name	Std.Dev.
plot	(Intercept)	3.35
Residual		3.04

number of obs: 1000, groups: plot, 10

AIC = 5116.3, DIC = 5096.3

deviance = 5102.3

lmer coefficients

```
coef(mixed)
```

```
$plot
```

	(Intercept)	dbh.c
1	32.13118	0.6169271
2	38.61479	0.6169271
3	36.46547	0.6169271
4	34.06404	0.6169271
5	35.75313	0.6169271
6	36.30517	0.6169271
7	32.04003	0.6169271
8	27.30620	0.6169271
9	37.27097	0.6169271
10	34.39546	0.6169271

```
attr(,"class")
```

```
[1] "coef.mer"
```

Bayesian varying-intercept model with pooling across plots

```
model3 <- function(){  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)    # tau = precision (inverse of variance)  
    mu[i] <- alfa[plot[i]] + beta*dbh[i]    # centred diameter  
  }  
  # PRIORS  
  for (j in 1:10){  
    alfa[j] ~ dnorm(grandmu, tauplot)    # Now we are estimating the plot variance!  
  }  
  grandmu ~ dnorm(0, .001)    # Overall mean height across all plots  
  tauplot <- pow(sigmaplot, -2)  
  sigmaplot ~ dunif(0, 20)    # between-plot variance  
  beta ~ dnorm(0, .001)  
  tau <- pow(sigma, -2)  
  sigma ~ dunif(0, 50)    # residual variance  
}
```

Call JAGS

```
data <- list(height=trees$height,  
             dbhc=trees$dbh.c,  
             plot=trees$plot)  
m3 <- jags(data,  
           model.file=model3,  
           parameters.to.save = c("alfa", "beta", "sigma", "grandmu", "sigmaplot"),  
           n.chains=3,  
           inits=NULL,  
           n.iter=10000,  
           n.burnin=5000)
```

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 4884

Initializing model

Results


Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpwpZd1g/model18fc4e2f632e.txt", f
 3 chains, each with 10000 iterations (first 5000 discarded), n.thin = 5
 n.sims = 3000 iterations saved

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alfa[1]	32.131	0.156	31.822	32.421	1.001	3000
alfa[2]	38.611	0.211	38.206	39.029	1.002	1400
alfa[3]	36.474	0.325	35.835	37.105	1.001	2100
alfa[4]	34.068	0.318	33.448	34.694	1.001	3000
alfa[5]	35.746	0.305	35.158	36.352	1.002	1500
alfa[6]	36.305	0.392	35.499	37.068	1.001	3000
alfa[7]	32.044	0.667	30.731	33.345	1.001	3000
alfa[8]	27.288	0.896	25.541	29.064	1.002	1400
alfa[9]	37.294	1.037	35.254	39.425	1.001	3000
alfa[10]	34.363	1.281	31.897	36.861	1.001	3000
beta	0.617	0.008	0.602	0.632	1.001	2700
grandmu	34.336	1.304	31.599	36.871	1.001	3000
sigma	3.048	0.071	2.916	3.186	1.001	3000
sigmaplot	3.942	1.237	2.304	6.982	1.001	3000
deviance	5064.988	5.105	5057.256	5076.492	1.002	1500


For each parameter, n.eff is a crude measure of effective sample size,
 and Rhat is the potential scale reduction factor (at convergence. Rhat=1).

48/73

A plot of the whole model

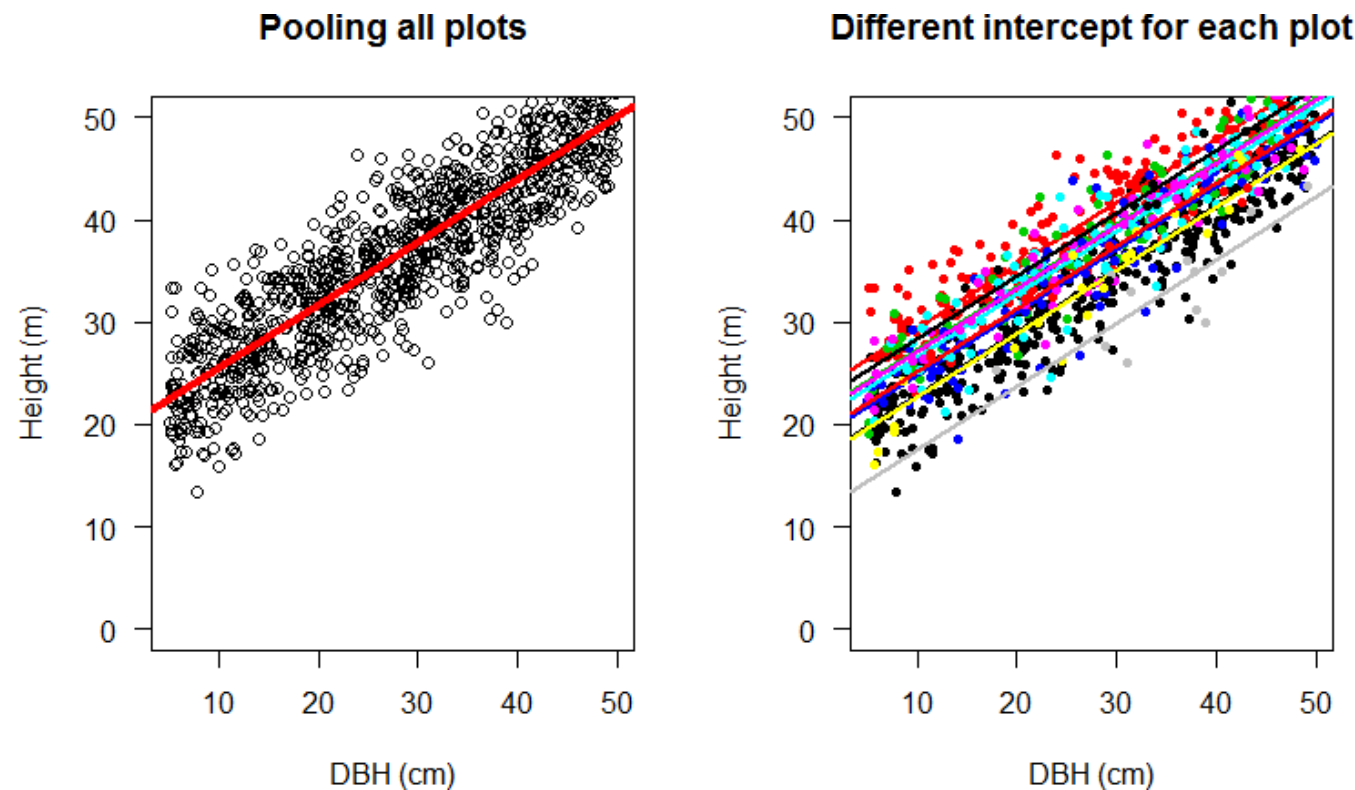


Comparing plot coefficients



A gradient from complete to no pooling

The multilevel model (with pooling) is somewhere between the complete-pooling (single intercept) and the no-pooling (one intercept for each plot, without shrinkage) models.



Growing the hierarchy: adding plot-level predictors

Model with group-level predictors

We had:

$$\begin{aligned}y_i &= a_j + bx_i + \varepsilon_i \\a_j &\sim N(0, \tau^2) \\\varepsilon_i &\sim N(0, \sigma^2)\end{aligned}$$

Now

$$\begin{aligned}y_i &= a_j + bx_i + \varepsilon_i \\a_j &\sim N(\mu_j, \tau^2) \\\mu_j &= \gamma + \delta \cdot \text{predictor}_j \\\varepsilon_i &\sim N(0, \sigma^2)\end{aligned}$$

Reading plot data

```
plotdata <- read.csv("plotdata.csv")  
temp.c <- plotdata$temp - 15
```

Model with group-level predictors

```
model4 <- function(){  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)  
    mu[i] <- alfa[plot[i]] + beta*dbhc[i]  
  }  
  # PRIORS  
  for (j in 1:10){  
    alfa[j] ~ dnorm(grandmu + beta.temp*tempc[j], tauplot)  
  }  
  beta.temp ~ dnorm(0, .001)  # slope for temperature effects  
  grandmu ~ dnorm(0, .001)  
  tauplot <- pow(sigmaplot, -2)  
  sigmaplot ~ dunif(0, 20)  
  beta ~ dnorm(0, .001)  
  tau <- pow(sigma, -2)  
  sigma ~ dunif(0, 50)  
}
```

running JAGS...

```
data <- list(height=trees$height,  
             dbhc=trees$dbh.c,  
             plot=trees$plot,  
             tempc=temp.c)  
m4 <- jags(data,  
           model.file=model4,  
           parameters.to.save = c("alfa", "beta", "sigma", "grandmu", "sigmaplot", "beta.temp"),  
           n.chains=3,  
           inits=NULL,  
           n.iter=10000,  
           n.burnin=5000)
```

Compiling model graph

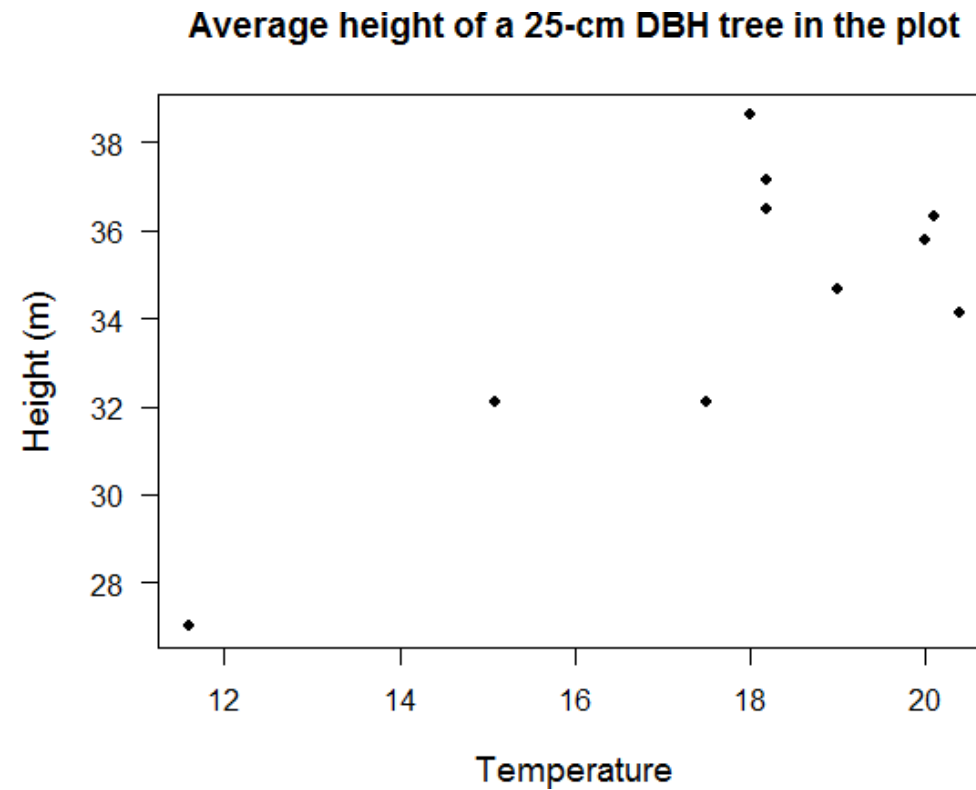
Resolving undeclared variables

Allocating nodes

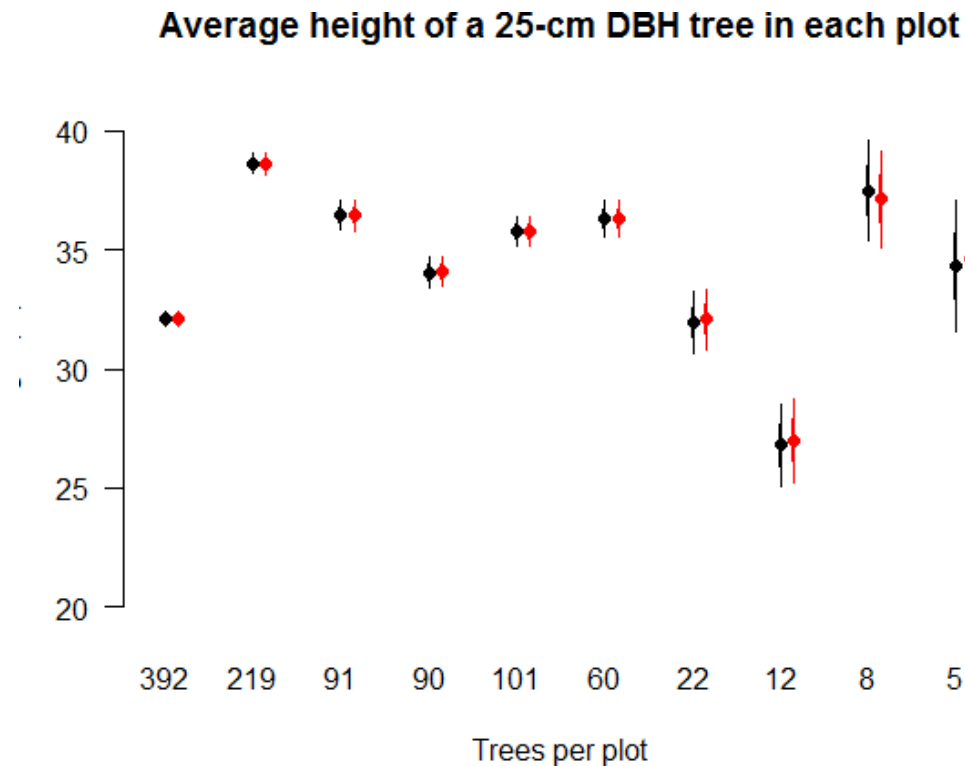
Graph Size: 4913

Initializing model


Average heights among plots related to temperature



Adding plot-level predictors (pooling) may improve parameter estimation



Adding plot-level predictors (pooling) may improve parameter estimation



From Gelman & Hill p. 253

Slopes can also vary...

- and coefficients be estimated with pooling
- but the correlation between slopes and intercepts must be modelled explicitly
- see e.g. Gelman & Hill 2007, ch. 13.

So what's a multilevel/hierarchical model?

Parameters/coefficients are given a probability model (with their own hyperparameters estimated from data).

Intercepts and/or slopes may vary, and can be modelled (sometimes including their own predictors).

Advantages of hierarchical Bayes


- Perfect for structured data (space-time)
- Predictors enter at the appropriate level
- Accommodate variation in treatment effects
- More efficient inference of regression parameters
- Using all the data to perform inferences for groups with small sample size
- Predictions fully accounting for uncertainty and variability
- Prior information

Datasets are stochastic realisations of a process


```
data1=rnorm(100, 2 + 1.6*x, 5)
```

```
data2=rnorm(100, 2 + 1.6*x, 5)
```

Histogram of data1

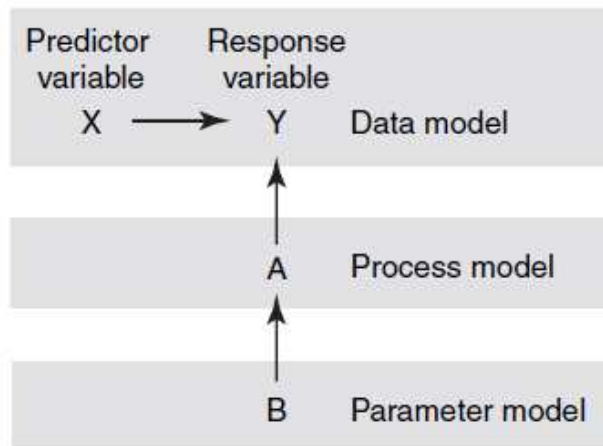


Histogram of data2



These two datasets are different, even though they arise from same process

Hierarchical Bayes: data, process, parameters



$$\begin{aligned} f(\text{data}, \text{process}, \text{parameters}) \\ \propto f(\text{data} | \text{process}, \text{parameters}) \\ \times f(\text{process} | \text{parameters}) \\ \times f(\text{parameters}). \end{aligned}$$

Clark et al. 2006, Clark 2007


Exercise

Does sex influence height?

Do it yourself

Bayesian logistic regression

Relationship between tree size and mortality



Logistic regression model

```
model5 <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(dead)){  
    dead[i] ~ dbern(pdeath[i])  
    logit(pdeath[i]) <- mu + beta*dbhc[i]  
  }  
  
  # PRIORS  
  mu ~ dnorm(0, .001)  
  beta ~ dnorm(0, .001)  
}
```

Calling JAGS

```
data <- list(dead=trees$dead,
             dbhc=trees$dbh.c)
m5 <- jags(data,
           model.file=model5,
           parameters.to.save = c("mu", "beta"),
           n.chains=3,
           inits=NULL,
           n.iter=10000,
           n.burnin=5000)
```

Compiling model graph

Resolving undeclared variables

Allocating nodes

Graph Size: 4668

Initializing model

Results

```
Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpwpZdlg/model18fc588f6cea.txt", f
  3 chains, each with 10000 iterations (first 5000 discarded), n.thin = 5
n.sims = 3000 iterations saved
```

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
beta	0.055	0.014	0.029	0.084	1.002	1600
mu	-3.461	0.222	-3.916	-3.076	1.002	1300
deviance	345.865	5.831	343.745	351.241	1.001	3000

For each parameter, n.eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor (at convergence, Rhat=1).

```
DIC info (using the rule, pD = var(deviance)/2)
```

```
pD = 17.0 and DIC = 362.9
```

DIC is an estimate of expected predictive error (lower deviance is better).

Compare with glm

```
logreg <- glm(dead ~ dbh.c, data = trees, family = binomial)
display(logreg)
```

```
glm(formula = dead ~ dbh.c, family = binomial, data = trees)
```

	coef.est	coef.se
(Intercept)	-3.44	0.21
dbh.c	0.05	0.01

```
n = 1000, k = 2
```

```
residual deviance = 343.7, null deviance = 360.9 (difference = 17.2)
```


END

